

S Y S D I G



W E B I N A R S

# Top Dockerfile Best Practices



**ALVARO IRADIER**

Integrations Engineer,  
Sysdig



# Agenda

- Why
- Shifting left security
- 1. Unnecessary privileges
- 2. Reduce attack surface
- 3. Credentials & Confidentiality
- 4. Linting & Scanning
- 5. Beyond image building

Examples available at: <https://github.com/airadier/webinar-dockerfile-best-practices>



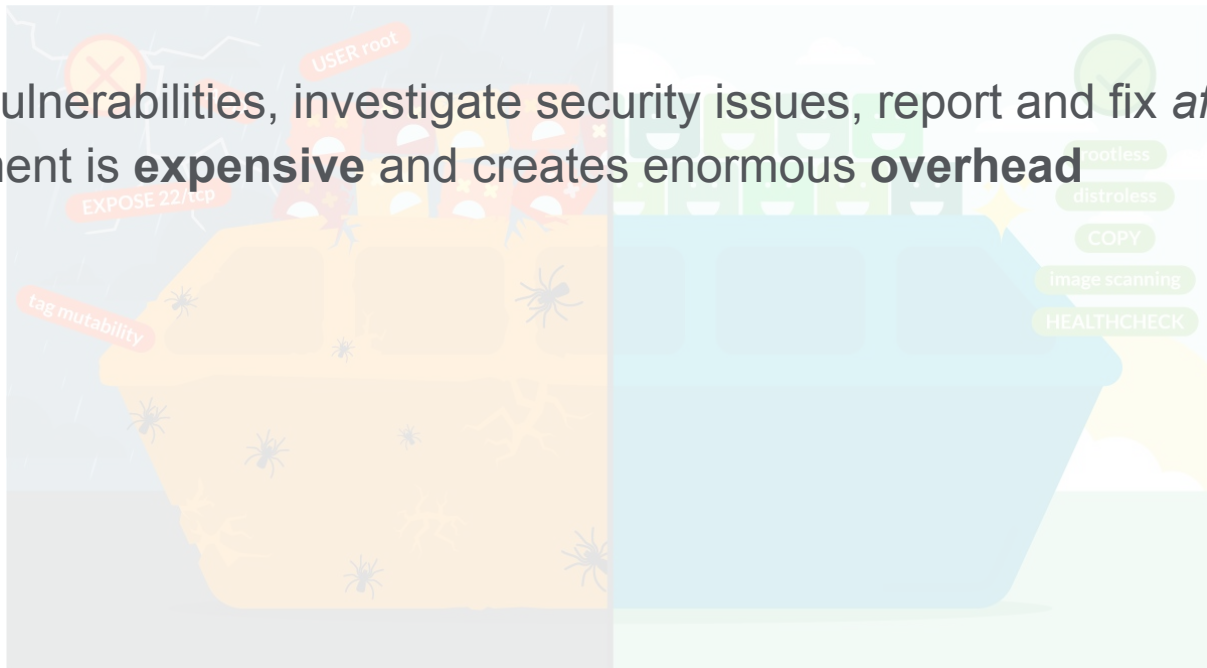
# Why

# Why



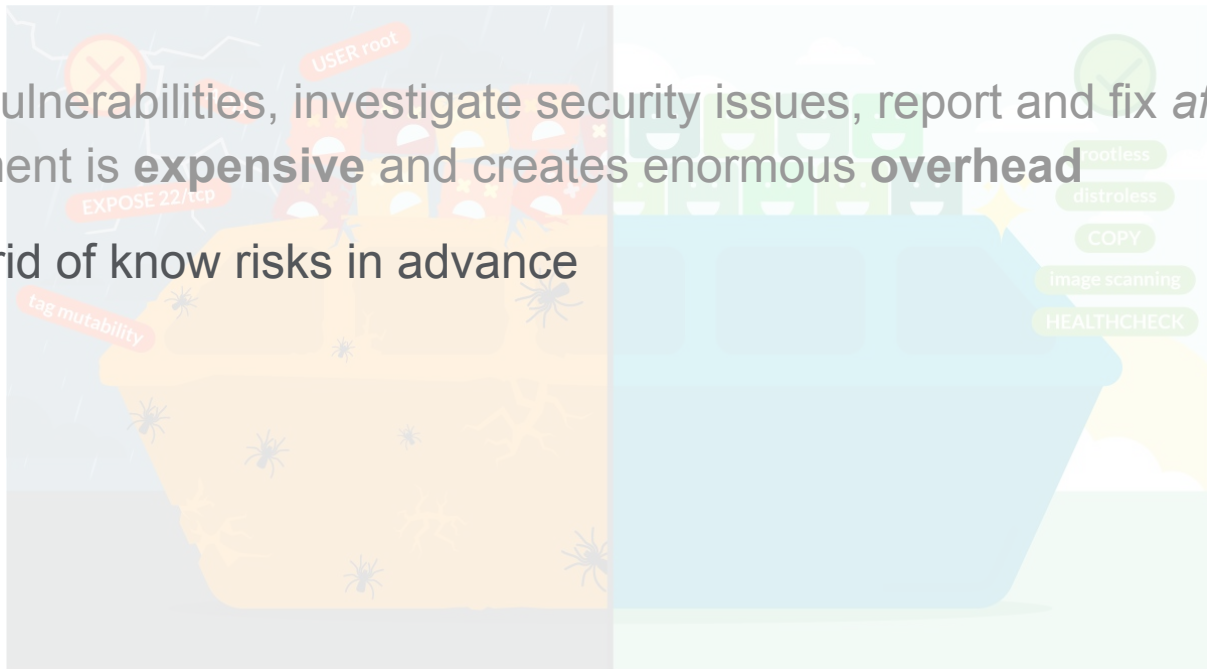
# Why

- Detect vulnerabilities, investigate security issues, report and fix *after* deployment is **expensive** and creates enormous **overhead**



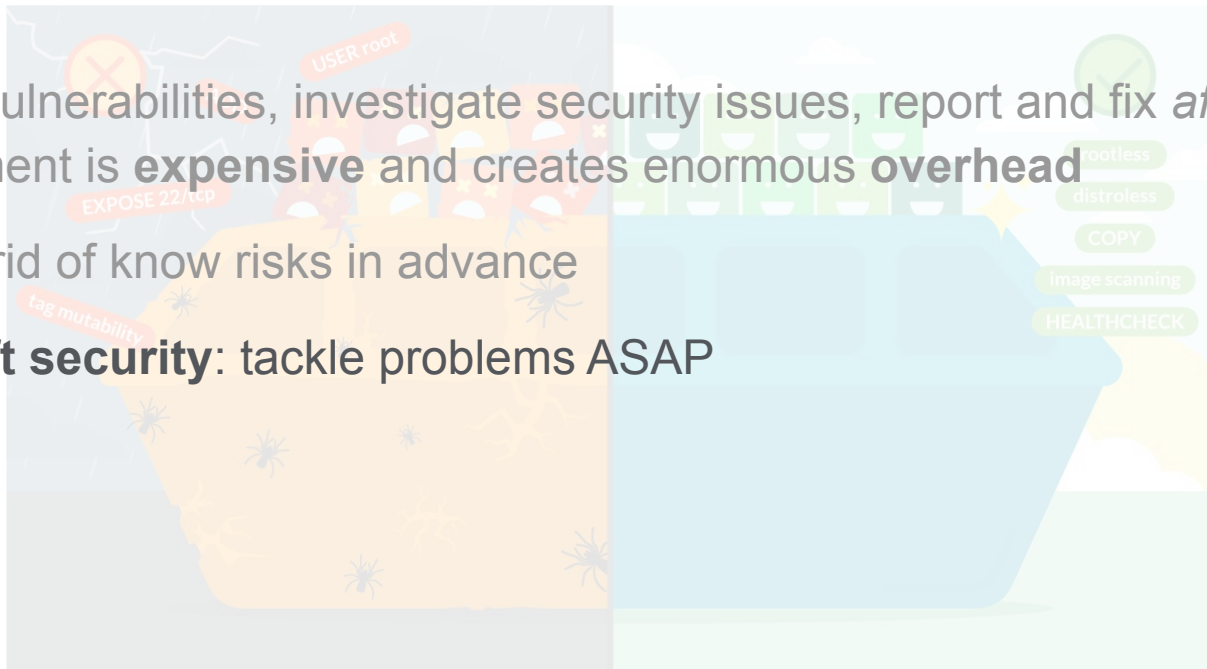
# Why

- Detect vulnerabilities, investigate security issues, report and fix *after* deployment is **expensive** and creates enormous **overhead**
- So, get rid of know risks in advance



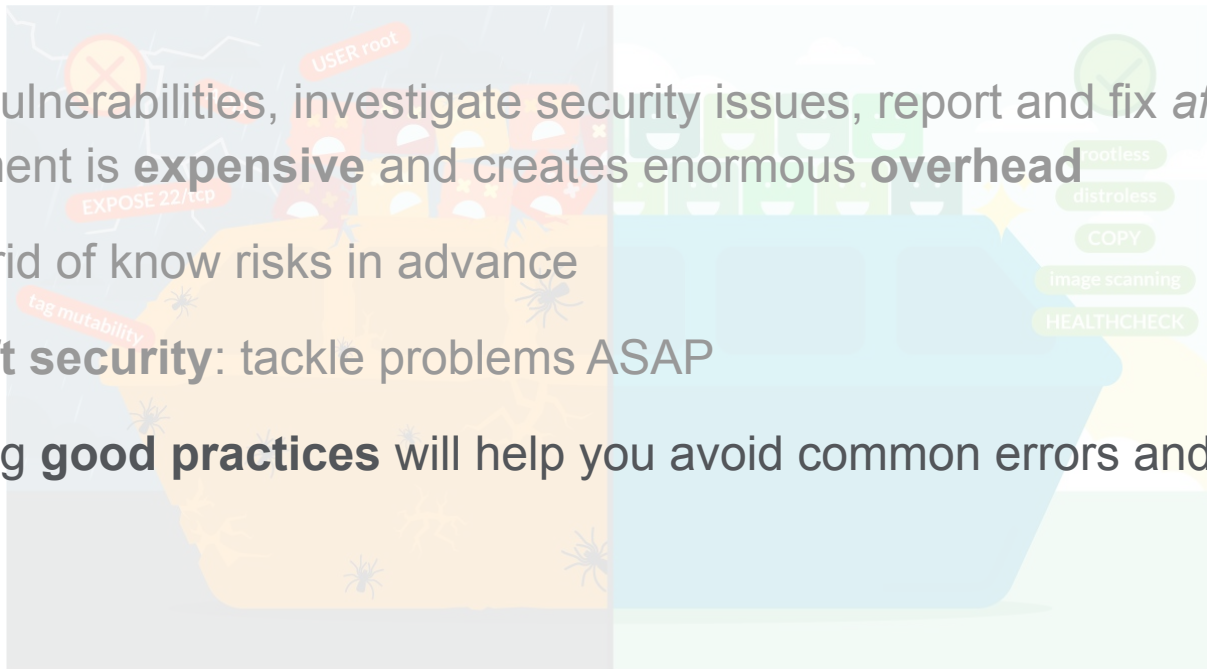
# Why

- Detect vulnerabilities, investigate security issues, report and fix *after* deployment is **expensive** and creates enormous **overhead**
- So, get rid of known risks in advance
- **Shift left security**: tackle problems ASAP



# Why

- Detect vulnerabilities, investigate security issues, report and fix *after* deployment is **expensive** and creates enormous **overhead**
- So, get rid of known risks in advance
- **Shift left security**: tackle problems ASAP
- Following **good practices** will help you avoid common errors and pitfalls







# Shifting left security

# Shifting left security

Security traditionally **managed later** on the software life cycle:

# Shifting left security

Security traditionally managed later on the software life cycle:

- You deliver your part of the cycle.

# Shifting left security

Security traditionally managed later on the software life cycle:

- You deliver your part of the cycle.
- Whatever happens next is “Somebody Else’s Problem”

## Somebody Else's Problem Field (S.E.P. Field)

The Somebody Else's Problem Field, or SEP, is a cheap easy and staggeringly useful way of protecting something from unwanted eyes.....Any object around which an SEP is applied will cease to be noticed, because any problems one may have become Somebody Else's.  
(Douglas Adams, Hitchhikers.wikia.com)



# Shifting left security

Security traditionally managed later on the software life cycle:

- You deliver your part of the cycle.
- Whatever happens next is "Somebody Else's Problem"
- Periodic auditor checks

## Somebody Else's Problem Field (S.E.P. Field)

The Somebody Else's Problem Field, or SEP, is a cheap easy and staggeringly useful way of protecting something from unwanted eyes.....Any object around which an SEP is applied will cease to be noticed, because any problems one may have become Somebody Else's.  
(Douglas Adams, Hitchhikers.wikia.com)



# Shifting left security

Security traditionally managed later on the software life cycle:

- You deliver your part of the cycle.
- Whatever happens next is “Somebody Else’s Problem”
- Periodic auditor checks
- InfoSec recommendations, benchmarks and validation: often seen by other parts of the company as a “barrier” that slows down the deployment of newest & shiniest software versions

## Somebody Else's Problem Field (S.E.P. Field)

The Somebody Else's Problem Field, or SEP, is a cheap easy and staggeringly useful way of protecting something from unwanted eyes.....Any object around which an SEP is applied will cease to be noticed, because any problems one may have become Somebody Else's.  
(Douglas Adams, Hitchhikers.wikia.com)



# Shifting left security

Security traditionally managed later on the software life cycle:

- You deliver your part of the cycle.
- Whatever happens next is “Somebody Else’s Problem”
- Periodic auditor checks
- InfoSec recommendations, benchmarks and validation: often seen by other parts of the company as a “barrier” that slows down the deployment of newest & shiniest software versions

**Shifting-left** means basically two things

## Somebody Else's Problem Field (S.E.P. Field)

The Somebody Else's Problem Field, or SEP, is a cheap easy and staggeringly useful way of protecting something from unwanted eyes.....Any object around which an SEP is applied will cease to be noticed, because any problems one may have become Somebody Else's.  
(Douglas Adams, Hitchhikers.wikia.com)



# Shifting left security

Security traditionally managed later on the software life cycle:

- You deliver your part of the cycle.
- Whatever happens next is “Somebody Else’s Problem”
- Periodic auditor checks
- InfoSec recommendations, benchmarks and validation: often seen by other parts of the company as a “barrier” that slows down the deployment of newest & shiniest software versions

**Shifting-left** means basically two things

- Implementing security closer to the earlier stages of the life cycle, into the **developers realm** (repositories and builds)

## Somebody Else's Problem Field (S.E.P. Field)

The Somebody Else's Problem Field, or SEP, is a cheap easy and staggeringly useful way of protecting something from unwanted eyes.....Any object around which an SEP is applied will cease to be noticed, because any problems one may have become Somebody Else's.  
(Douglas Adams, Hitchhikers.wikia.com)





# Shifting left security

Security traditionally managed later on the software life cycle:

- You deliver your part of the cycle.
- Whatever happens next is “Somebody Else’s Problem”
- Periodic auditor checks
- InfoSec recommendations, benchmarks and validation: often seen by other parts of the company as a “barrier” that slows down the deployment of newest & shiniest software versions

**Shifting-left** means basically two things

- Implementing security closer to the earlier stages of the life cycle, into the **developers realm** (repositories and builds)
- **Everyone shares the security responsibilities**, and collaborates towards common goals

## Somebody Else's Problem Field (S.E.P. Field)

The Somebody Else's Problem Field, or SEP, is a cheap easy and staggeringly useful way of protecting something from unwanted eyes.....Any object around which an SEP is applied will cease to be noticed, because any problems one may have become Somebody Else's.  
(Douglas Adams, Hitchhikers.wikia.com)



# Shifting left security

Security traditionally managed later on the software life cycle:

- You deliver your part of the cycle.
- Whatever happens next is “Somebody Else’s Problem”
- Periodic auditor checks
- InfoSec recommendations, benchmarks and validation: often seen by other parts of the company as a “barrier” that slows down the deployment of newest & shiniest software versions

**Shifting-left** means basically two things

- Implementing security closer to the earlier stages of the life cycle, into the **developers realm** (repositories and builds)
- **Everyone shares the security responsibilities**, and collaborates towards common goals

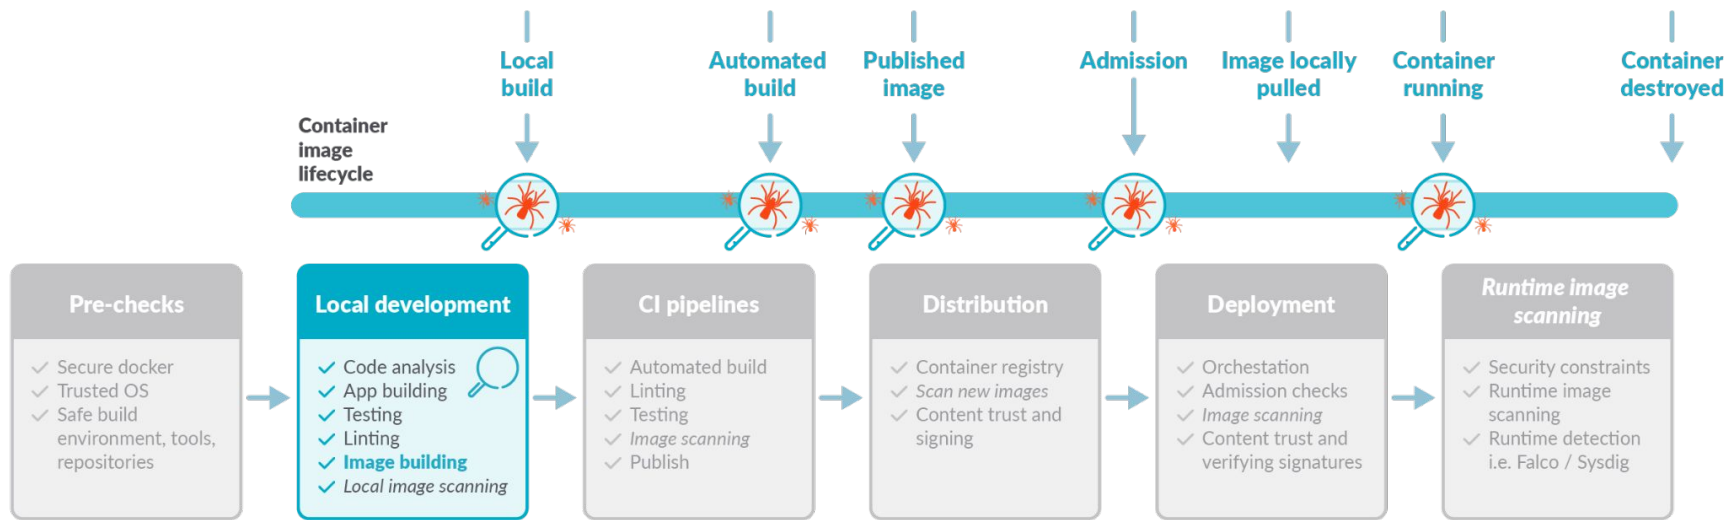
**Dockerfile best practices** are applied in **earlier stages**

## Somebody Else's Problem Field (S.E.P. Field)

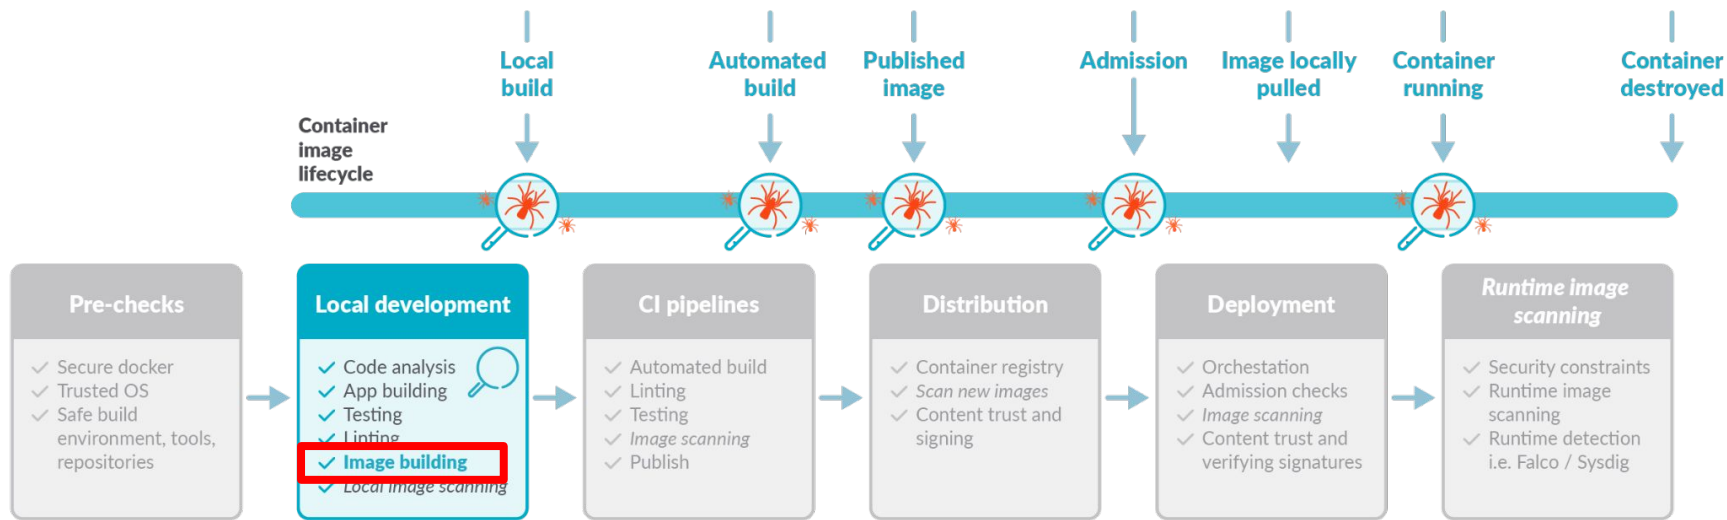
The Somebody Else's Problem Field, or SEP, is a cheap easy and staggeringly useful way of protecting something from unwanted eyes.....Any object around which an SEP is applied will cease to be noticed, because any problems one may have become Somebody Else's.  
(Douglas Adams, Hitchhikers.wikia.com)



# Shifting left security



# Shifting left security






# 1. Unnecessary privileges

# Unnecessary privileges

- **Problem:** containers often running with more privileges than required
  - Our [recent report](#) highlighted that 58% of images are running the container entrypoint as root (UID 0)



- **Risks:**
  - Attackers can exploit vulnerabilities or bugs to gain access to other services or resources
  - <https://sysdig.com/blog/lateral-movement-cloud-containers/>
  - Openshift or others blocking root containers



# Unnecessary privileges

- Step 1: Initial Dockerfile

```
FROM alpine
COPY ./src/example_app /example_app
ENTRYPOINT /example_app
```

```
> docker build -f Dockerfile1 -t test1_1 .
[+] Building 0.5s (7/7) FINISHED

> docker run --name test --rm -p 5000:5000 test1_1
Listening at :5000
```

# Unnecessary privileges

- Testing the app

```
> curl "localhost:5000/?name=webinar"
Hi webinar!
> curl "localhost:5000/?name=alvaro"
Hi alvaro!
> docker exec -ti test sh
/ # ls -lh
-rw-r--r--      1 root      root           111 Apr  7 11:36 access.log
-rwxr-xr-x      1 root      root        5.9M Apr  7 11:27 example_app
...
/ # cat access.log
[2021-04-07T11:36:14Z] IP=172.17.0.1:60974 name=webinar
[2021-04-07T11:36:16Z] IP=172.17.0.1:60978 name=alvaro
/ # ps
PID     USER      TIME  COMMAND
   1   root          0:00  /example_app
```



# Unnecessary privileges

- Step 2: running as non-root

```
FROM alpine
COPY ./src/example_app /app/example_app
WORKDIR /app
USER 1000
ENTRYPOINT /app/example_app
```

```
> docker build -f Dockerfile2 -t test1_2 .
[+] Building 0.5s (7/7) FINISHED
> docker run --name test --rm -p 5000:5000 test1_2
panic: open access.log: permission denied
```

# Unnecessary privileges

- Step 3: fix permissions

```
FROM alpine
COPY ./src/example_app /app/example_app
WORKDIR /app
RUN chown 1000:1000 /app
USER 1000
ENTRYPOINT /app/example_app
```

```
> docker build -f Dockerfile3 -t test1_3 .
[+] Building 0.5s (7/7) FINISHED
> docker run --name test --rm -p 5000:5000 test1_3
Listening at :5000
```

# Unnecessary privileges

- Testing the **rootless** app

```
...
> docker exec -ti test sh
/app $ cat access.log
[2021-04-07T11:45:56Z] IP=172.17.0.1:60996 name=webinar
[2021-04-07T11:45:58Z] IP=172.17.0.1:61000 name=alvaro
/app $ ls -lh
total 6M
-rw-r--r--      1 1000      root           111 Apr  7 11:45 access.log
-rwxr-xr-x      1 root      root           5.9M Apr  7 11:27 example_app
/app $ ps
PID   USER      TIME  COMMAND
   1   1000      0:00  /app/example_app
```

# Unnecessary privileges

- Step 4: unspecific UID

```
> docker run -u 1001 --name test --rm -p 5000:5000 test1_3  
panic: open access.log: permission denied
```

```
FROM alpine  
COPY ./src/example_app /app/example_app  
USER 1000  
WORKDIR /tmp  
ENTRYPOINT /app/example_app
```

```
> docker run -u 1001 --name test --rm -p 5000:5000 test1_4  
Listening at :5000
```

# Unnecessary privileges

- Verifying the container is running with UID 1001

```
...  
> docker exec -ti test sh  
/tmp $ ls -lh  
total 4K  
-rw-r--r--    1 1001      root          111 Apr  7 11:53 access.log  
/tmp $ ps  
PID   USER      TIME  COMMAND  
    1  1001      0:00  /app/example_app
```

# Unnecessary privileges

- **Recap:**
  - Step1: We started from a container running with default root user
  - Step 2: We make it run as **non-root** user, and find a permissions issue
  - Step 3: We fix permissions in the Dockerfile
    - Try to run it as other user UID, permissions issue again
  - Step 4: Allow running as **any user** by writing to /tmp
    - Separate app folder and data folder - we can have **persistence**

# Unnecessary privileges

- **Prevention:**
  - Follow the principle of **least privilege** so your service or application only has access to the resources and information necessary to perform its purpose
  - **USER directive**, run as non-root by default
  - Allow running with **random UIDs**
    - Required in some environments, as Openshift
    - Simplifies permissions with host mounts: match container and host UIDs



## 2. Reduce attack surface



# Reduce attack surface

- **Problem:** including unnecessary packages or exposing unused ports
- **Risks:**
  - Your system is more exposed to attacks
  - Using components not under your control

# Reduce attack surface

- **Don't:**
  - Build the application externally, copy into the container
  - Bad reproducibility

```
FROM alpine
COPY ./src/example_app /example_app
USER 1000
WORKDIR /tmp
ENTRYPOINT /example_app
```

# Reduce attack surface

- **Don't:**
  - **Build directly inside the final container**
    - Big image size and multiple layers, build toolchain included in the container, other unrequired packages, remainings of the application source code, ...

```
FROM alpine
RUN apk add go
COPY ./src/ /src
WORKDIR /src
RUN go build .
ENTRYPOINT /src/example_app
```

# Reduce attack surface

- **Multistage builds:**

```
#This is the "builder" stage
```

```
FROM golang:1.16 as builder
```

```
WORKDIR /my-go-app
```

```
COPY src .
```

```
RUN GOOS=linux GOARCH=amd64 go build .
```

```
#This is the final stage, and we copy artifacts from "builder"
```

```
FROM alpine
```

```
COPY --from=builder /my-go-app/example_app /bin/example_app
```

```
ENTRYPOINT ["/bin/example_app"]
```

# Reduce attack surface

- **Multistage builds:**
  - Reproducible builds, always same build environment
  - Minimal image size, no build tools or undesired packages

```
> docker images
test2_1          latest          681b9e590ae9   5 minutes ago   448MB
test2_2          latest          37d02efde1e4   35 seconds ago  7.55MB
...
```

# Reduce attack surface

- **Example: multistage build with nodejs and typescript**

(<https://github.com/kevinpollet/typescript-docker-multi-stage-build>):

```
FROM node:14-alpine AS builder
WORKDIR /usr/src/app
COPY typescript-docker-multi-stage-build/package*.json ./
RUN npm ci
COPY typescript-docker-multi-stage-build/tsconfig*.json ./
COPY typescript-docker-multi-stage-build/src src
RUN npm run build
```

```
FROM node:14-alpine
ENV NODE_ENV=production
WORKDIR /usr/src/app
RUN chown node:node .
USER node
COPY typescript-docker-multi-stage-build/package*.json ./
RUN npm install
COPY --from=builder /usr/src/app/lib/ lib/
EXPOSE 3000
ENTRYPOINT [ "node", "lib/server.js" ]
```

# Reduce attack surface

- **Don't:** Use big, generic distro images if not needed (i.e. ubuntu)

```
> docker run ... quay.io/sysdig/secure-inline-scan:2 image-ubuntu -k $SYSDIG_SECURE_TOKEN --storage-type docker-daemon
Inspecting image from Docker daemon -- distroless-1:latest
Full image: docker.io/library/image-ubuntu
Full tag: docker.io/library/image-ubuntu:latest
...
Analyzing image..
Analysis complete!
...
Evaluation results
- warn dockerfile:instruction Dockerfile directive 'HEALTHCHECK' not found, matching condition 'not_exists' check
- warn dockerfile:instruction Dockerfile directive 'USER' not found, matching condition 'not_exists' check
- warn files:suid_or_guid_set SUID or SGID found set on file /bin/mount. Mode: 0o104755
- warn files:suid_or_guid_set SUID or SGID found set on file /bin/su. Mode: 0o104755
- warn files:suid_or_guid_set SUID or SGID found set on file /usr/bin/chage. Mode: 0o102755
...
Vulnerabilities report
Vulnerability Severity Package Type Fix version URL
- CVE-2019-18276 Low bash-4.3-14ubuntu1.4 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2019-18276
- CVE-2016-2781 Low coreutils-8.25-2ubuntu3~16.04 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-2781
- CVE-2017-8283 Negligible dpkg-1.18.4ubuntu1.6 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2017-8283
- CVE-2020-13844 Medium gcc-5-base-5.4.0-6ubuntu1~16.04.12 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2020-13844
...
- CVE-2018-20839 Medium systemd-sysv-229-4ubuntu21.29 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-20839
- CVE-2016-5011 Low util-linux-2.27.1-6ubuntu3.10 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-5011
...
```

# Reduce attack surface

- **Don't:** Use big, generic distro images if not needed (i.e. ubuntu)

- Things most likely won't need in your final image:

- gcc-5 compiler

- sysV compatibility

- dpkg? bash?


- More than 100 vulnerabilities detected!

```
> docker run -v /var/run/docker.sock:/var/run/docker.sock --rm quay.io/sysdig/secure-inline-scan:2 image-ubuntu -k $SYSDIG_SECURE_TOKEN --storage-type
docker-daemon
Importing image from Docker daemon -- distroless:latest
Full tarball localbuild/distroless:latest
...
Analyzing image...
Analysis completed!
...
Ev...
Vulnerabilities report
Vulnerability Severity Package Type Fix version URL
- CVE-2016-2781 Low coreutils-8.25-ubuntu3~16.04 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-2781
- CVE-2017-8283 Negligible dpkg-1.18.4ubuntu1.6 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2017-8283
- CVE-2020-13844 Medium gcc-5-base-5.4.0-6ubuntu1~16.04.12 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2020-13844
...
- CVE-2018-20839 Medium systemd-sysv-229-4ubuntu21.29 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2018-20839
- CVE-2016-5011 Low util-linux-2.27.1-6ubuntu3.10 dpkg None http://people.ubuntu.com/~ubuntu-security/cve/CVE-2016-5011
...
```



# Reduce attack surface

- **Don't: Official images might not be the best fit *per se***
  - Regarding security and minimalism, they might not be updated that often and can include extra packages for general use cases
- **Don't: Use outdated images**
  - New security vulnerabilities are discovered continuously
  - Stick to the latest security patches
  - No need to always go with the latest version (breaking changes)



# Reduce attack surface

- **Don't:** Use `docker.io/johndoe hacker/mycustom-node-image:latest`

```
FROM docker.io/johndoe hacker/mycustom-node-image:latest  
...
```

- Inherit all of the problems and vulnerabilities from that image
- Who builds and publishes that image?
- Is it updated regularly?
- How is it built?
- Are we sure the published version is really from the public Dockerfile?

# Reduce attack surface

- Prefer **verified** and **official** images from **trusted repositories** and **providers**
- Check for optimized vs generic versions
  - **Example:** [bitnami/node](#) vs [official node image](#)
    - customized versions on top of a minideb distribution
    - frequently updated with the latest bug fixes
    - signed with Docker Content Trust
    - pass a [security scan for tracking known vulnerabilities](#)

# Reduce attack surface

- **Be minimal**

- *alpine* versions
- FROM scratch
- Distroless (<https://github.com/GoogleContainerTools/distroless>):
  - i.e.: FROM `gcr.io/distroless/base-debian10`
    - Basic set of packages, including just required libraries like *glibc*, *libssl*, and *openssl*.
  - Slimmer: FROM `gcr.io/distroless/static-debian10`
    - For statically compiled applications like Go that don't require libc

# Reduce attack surface



# Reduce attack surface

- When using custom images, check for the **image source** and the Dockerfile, and **build your own base image**
- Define a **versioning strategy**:
  - Stick to stable or long-term support versions
  - Rebuild periodically
    - To get the latest packages from the base distro
    - npm, go mod offer ways to specify version ranges (keep up with latest security updates)
  - Plan in advance.
    - Be ready to drop old versions
    - Migrate before base image reaches the end of its life and stops receiving updates

# Reduce attack surface

- **Prevention:**

- Keep images *minimal*, only required stuff should be included
- Carefully choose the image that best fits your use case
- Use trusted, verified base images
- Use stable and well maintained versions, with frequent updates
- Update and rebuild your own images often



## 3. Credentials & Confidentiality



# Credentials & Confidentiality

- **Problem:** leaking credentials or confidential information in your images
- **Risks:**
  - Attackers can use leaked credentials to access your systems
  - Exposure of confidential or sensitive information

# Credentials & Confidentiality

- **Don't:**
  - Include hard coded credentials
  - Add credentials file or environment variables

```
FROM alpine
...
ENV SECURE_API_TOKEN=ajhda8-12312-29889234-foo
COPY aws_credentials /home/app/.aws/credentials
...
ENTRYPOINT /example_app
```

# Credentials & Confidentiality

```
> docker inspect test3_1
[
  {
    "Id": "sha256:18440a2433ea49efa686febc6f02c21a652a498523ed42e00cf79ebf3717cc0a",
    "RepoTags": [
      "test3_1:latest"
    ],
    "RepoDigests": [],
    "Parent": "",
    "Comment": "buildkit.dockerfile.v0",
    "Created": "2021-04-07T17:01:00.8254965Z",
    "Container": "",
    "ContainerConfig": {
      ...
    },
    "DockerVersion": "",
    "Author": "",
    "Config": {
      ...
      "Env": [
        "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin",
        "SECURE_API_TOKEN=ajhda8-12312-29889234-foo"
      ]
    }
  }
]
```

# Credentials & Confidentiality

```
> docker run --entrypoint /bin/sh --rm test3_1 -c "cat /home/app/.aws/credentials"
[default]
aws_access_key_id = SOME-ACCESS-KEY
aws_secret_access_key = SOME-SECRET-KEY
```

# Credentials & Confidentiality

- **Even if the file is removed!!**

```
FROM alpine
...
COPY aws_credentials /home/app/.aws/credentials
...
RUN rm /home/app/.aws/credentials
ENTRYPOINT /example_app
```

# Credentials & Confidentiality

```
> docker run --entrypoint /bin/sh --rm test3_2 -c "cat /home/app/.aws/credentials"
cat: can't open '/home/app/.aws/credentials': No such file or directory
>
```

# Credentials & Confidentiality

```
> docker run --entrypoint /bin/sh --rm test3_2 -c "cat /home/app/.aws/credentials"
cat: can't open '/home/app/.aws/credentials': No such file or directory

> skopeo copy docker-daemon:test3_2:latest oci:test3_2
Getting image source signatures
Copying blob 8ea3b23f387b done
Copying blob 35c29c7d6159 done
Copying blob 6ddf15f6fc2b done
Copying config 6a1057f9fe done
Writing manifest to image destination
Storing signatures

>
```

# Credentials & Confidentiality

```
> docker run --entrypoint /bin/sh --rm test3_2 -c "cat /home/app/.aws/credentials"
cat: can't open '/home/app/.aws/credentials': No such file or directory

> skopeo copy docker-daemon:test3_2:latest oci:test3_2
Getting image source signatures
Copying blob 8ea3b23f387b done
Copying blob 35c29c7d6159 done
Copying blob 6ddf15f6fc2b done
Copying config 6a1057f9fe done
Writing manifest to image destination
Storing signatures

> cat test3_2/index.json
{"schemaVersion":2,"manifests":[{"mediaType":"application/vnd.oci.image.manifest.v1+json","digest":"sha256:a80e7da14ffc58f2d4f7e22ed6f71aaaa318a4ee8c605bbd47ea48f8ef5e9089","size":657}]}
```



# Credentials & Confidentiality

```
> docker run --entrypoint /bin/sh --rm test3_2 -c "cat /home/app/.aws/credentials"
cat: can't open '/home/app/.aws/credentials': No such file or directory

> skopeo copy docker-daemon:test3_2:latest oci:test3_2
Getting image source signatures
Copying blob 8ea3b23f387b done
Copying blob 35c29c7d6159 done
Copying blob 6ddf15f6fc2b done
Copying config 6a1057f9fe done
Writing manifest to image destination
Storing signatures

> cat test3_2/index.json
{"schemaVersion":2,"manifests":[{"mediaType":"application/vnd.oci.image.manifest.v1+json","digest":"sha256:a80e7da14ffc58f2d4f7e22ed6f71aaaa318a4ee8c605bbd47ea48f8ef5e9089","size":657}]}

> cat test3_2/blobs/sha256/a80e7da14ffc58f2d4f7e22ed6f71aaaa318a4ee8c605bbd47ea48f8ef5e9089 | jq
{
  "schemaVersion": 2,
  "config": {
    "mediaType": "application/vnd.oci.image.config.v1+json",
    "digest": "sha256:6a1057f9fe2693956a5bb40bd9e3ee624171f4145dc5b6c7d83b03e4d2774688",
    "size": 1236
  },
  "layers": [
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:2d5a20755f17e53a78fdfeebffff1100a88ec7941c727a9538932b0409ca7bf5c",
      "size": 2899855
    },
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:f4dca0e2aa585f2df801ce346c78580075396b610b1195887b11e74dbc4861f79"
```

# Credentials & Confidentiality

```
> cat test3_2/blobs/sha256/a80e7da14ffc58f2d4f7e22ed6f71aaaa318a4ee8c605bbd47ea48f8ef5e9089 | jq
{
  "schemaVersion": 2,
  "config": {
    "mediaType": "application/vnd.oci.image.config.v1+json",
    "digest": "sha256:6a1057f9fe2693956a5bb40bd9e3ee624171f4145dc5b6c7d83b03e4d2774688",
    "size": 1236
  },
  "layers": [
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:2d5a20755f17e53a78fdfeebfff1100a88ec7941c727a9538932b0409ca7bf5c",
      "size": 2899855
    },
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:f4dca0e2aa585f2df801ce346c78580075396b610b1195887b11e74dbc4861f7",
      "size": 284
    },
    {
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
      "digest": "sha256:27bc864279bbb3afc3ef38c2acd34b6c2897249e55076e518e7abc36f9ec4ec3",
      "size": 199
    }
  ]
}
```

# Credentials & Confidentiality

```
> tar xvfz test3_2/blobs/sha256/f4dca0e2aa585f2df801ce346c78580075396b610b1195887b11e74dbc4861f7
x home/
x home/app/
x home/app/.wh..wh..opq
x home/app/.aws/
x home/app/.aws/credentials

>
```

# Credentials & Confidentiality

```
> tar xvfz test3_2/blobs/sha256/f4dca0e2aa585f2df801ce346c78580075396b610b1195887b11e74dbc4861f7
x home/
x home/app/
x home/app/.wh..wh..opq
x home/app/.aws/
x home/app/.aws/credentials

> cat home/app/.aws/credentials
[default]
aws_access_key_id = SOME-ACCESS-KEY
aws_secret_access_key = SOME-SECRET-KEY
```

# Credentials & Confidentiality

- **Don't:**
  - Forget the **layered** nature of images (each command creates a new layer)
    - Removing a file in a layer still takes space and file can be accessed
    - Combine commands to reduce number of layers, i.e.:  
`RUN apt-get install wget && wget https://.../downloadedfile.tar && tar xvfz downloadedfile.tar && rm downloadedfile.tar && apt-get remove wget`
    - Optimize layers, place commands **less likely to change** (and easier to cache) **first**

# Credentials & Confidentiality

- **Example (unoptimized layers)**

```
FROM ubuntu
COPY source/* .
RUN apt-get update
RUN apt-get install -y wget nodejs
RUN wget https://bit.ly/3urGNtE -O downloadedfile.tgz
RUN tar xvzf downloadedfile.tgz
RUN rm downloadedfile.tgz
RUN apt-get -y remove wget
ENTRYPOINT ["/usr/bin/node", "/main.js"]
```

# Credentials & Confidentiality

- **Example (layer optimization)**

```
FROM ubuntu
COPY source/* .
RUN apt-get update && \
    apt-get install -y wget nodejs && \
    wget https://bit.ly/3urGNtE -O downloadedfile.tgz && \
    tar xvzf downloadedfile.tgz && \
    rm downloadedfile.tgz && \
    apt-get -y remove wget
ENTRYPOINT ["/usr/bin/node", "/main.js"]
```

# Credentials & Confidentiality

```
> docker images | grep test3
```

test3_3	latest	b022cc12df25	42 seconds ago	170MB
test3_4	latest	28ecf7863e57	16 seconds ago	168MB



# Credentials & Confidentiality

```
> docker inspect test3_3
[
  {
    "RootFS": {
      "Type": "layers",
      "Layers": [
        "sha256:0e64bafdc7ee828d0f3995bebfa388ced52a625ad2969eeb569f4a83db56d505",
        "sha256:935f303ebf75656fcbf822491f56646c5a875bd0ad0bf2529671d31dd5456dfa",
        "sha256:346be19f13b0ccad355ab89265edaa4ac5958a42b8bb0492d2d22d9e4538def4",
        "sha256:2a833093776af60022e7650aaec22cf1d6ef3a3aa6fb1dda965c79799e2af727",
        "sha256:e807acdc370d757d2a750a383b3f2486498a0188b786c79afae20557cc2e5145",
        "sha256:5d327508b8d97edd052a27133f01b69dd528912b086d21848dd0255fda647e3b",
        "sha256:faa1d0007907eec40fff21339c1f1b627ade08e542f990f47405dfb391837369",
        "sha256:9bee3fb1911cbf5d5f3b8c3d3eea60c494ad8443330621c26f656313a0d2ff37",
        "sha256:e360d699d4864a9ee3e5ed3623696b1b011a0146f103a197adf7b17926cb09a7",
        "sha256:b7359736dd9cc0981e91a037b0eb8061b92dcc2466000f30bf145998d02b48b3"
      ]
    },
    "Metadata": {
      "LastTagTime": "2021-04-08T11:50:15.6160675Z"
    }
  }
]
```

# Reduce attack surface

```
> docker inspect test3_4
[
  {
    "RootFS": {
      "Type": "layers",
      "Layers": [
        "sha256:0e64bafdc7ee828d0f3995bebf388ced52a625ad2969eeb569f4a83db56d505",
        "sha256:935f303ebf75656fcbf822491f56646c5a875bd0ad0bf2529671d31dd5456dfa",
        "sha256:346be19f13b0ccad355ab89265edaa4ac5958a42b8bb0492d2d22d9e4538def4",
        "sha256:2a833093776af60022e7650aaec22cf1d6ef3a3aa6fb1dda965c79799e2af727",
        "sha256:2854af5a66a664872a1a394b55ac85511d72ecb9d857238fd717feb59c3d963d"
      ]
    },
    "Metadata": {
      "LastTagTime": "2021-04-08T11:50:41.9055789Z"
    }
  }
]
```

# Credentials & Confidentiality

- **Example (cache optimization)**

```
FROM ubuntu
COPY source/* .
RUN apt-get update && \
    apt-get install -y wget nodejs && \
    wget https://bit.ly/3urGNtE -O downloadedfile.tgz && \
    tar xvzf downloadedfile.tgz && \
    rm downloadedfile.tgz && \
    apt-get -y remove wget
COPY source/* .
ENTRYPOINT ["/usr/bin/node", "/main.js"]
```

# Credentials & Confidentiality

- **Don't:** Leak files from the build context

```
docker build -t myimage .
```

- The “.” parameter is the build context
- All the files in the build context are sent to the docker daemon
  - You can copy confidential or unnecessary files into the container, like configuration files, credentials, backups, lock files, temporary files, sources, subfolders, dotfiles, etc.
- COPY and ADD commands work from the build context

# Credentials & Confidentiality

- **Example**

```
docker build -f Dockerfile -t myimage .
```

```
...  
COPY . /my-app
```

This would copy everything inside the build context, which for the “.” example, includes the Dockerfile itself.

# Credentials & Confidentiality

Good practices:

- Use a **clean build context**

```
docker build -f Dockerfile -t myimage files/
```

- Use **.dockerignore** file
- Prefer **COPY over ADD**, and **avoid wildcards**
  - COPY is more explicit, more predictable and less error prone
  - ADD can add files from a URL or from a .tar file



## 4. Linting & Scanning

# Linting

- Tools like [Haskell Dockerfile Linter](#) (hadolint) can detect bad practices in your Dockerfile, and even expose issues inside the shell commands executed by the RUN instruction.
- Image scanners (like Sysdig's) are also capable of detecting bad practices via customizable rules
- Automate:** Consider incorporating such a tool in your CI pipelines.

IMAGE SCANNING  
Policies > Edit Policy

Default policies are read-only. Duplicate this policy to make changes.

Name: Default Configuration Policy - Dockerfile Best Practices

Description: This policy provides out of the box rules around Dockerfile best practices. We frequently update these policies and if you'd like to modify the policy you should use this as a base template to avoid modifications being overwritten.

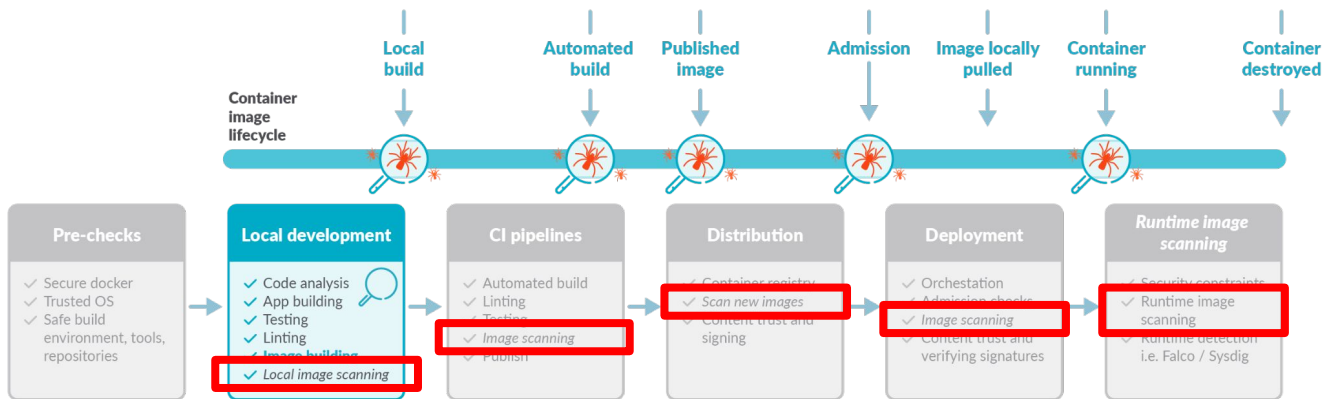
Rules				
Vulnerabilities	Stale feed data	Max days since sync: 7	Warn	X
Dockerfile	Instruction	Instruction: RUN; Check like: Value: *apt-get upgrade.*	Warn	X
Dockerfile	Instruction	Instruction: RUN; Check like: Value: *yum upgrade.*	Warn	X
Dockerfile	Instruction	Instruction: HEALTHCHECK; Check: not_exists	Warn	X
Dockerfile	Effective user	Type: blacklist; Users: root	Warn	X
Dockerfile	Exposed ports	Type: blacklist; Ports: 22	Warn	X
Dockerfile	Instruction	Instruction: LABEL; Check: !=; Value: latest	Warn	X
Dockerfile	Instruction	Instruction: ENV; Check like: Value: *passwordPASSWORDpasswordPASSWORDSecretSECRET*	Warn	X
Dockerfile	Instruction	Instruction: USER; Check: not_exists	Warn	X
Dockerfile	Instruction	Instruction: ADD; Check: exists	Warn	X

Dockerfile	Effective user	Type: blacklist; Users: root
Dockerfile	Exposed ports	Type: blacklist; Ports: 22

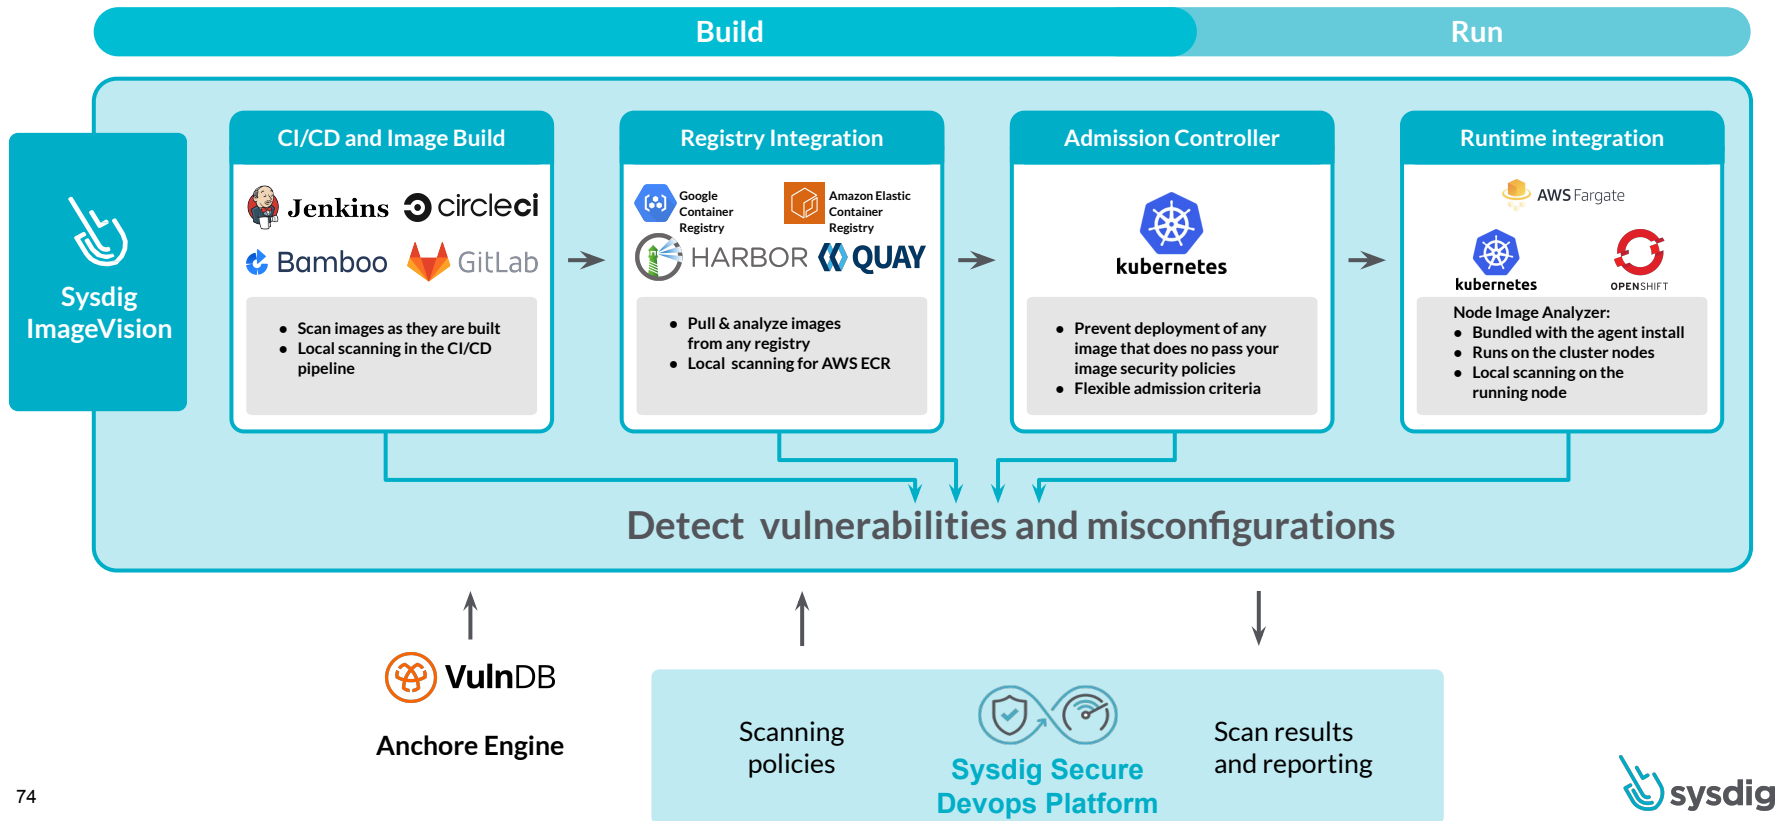


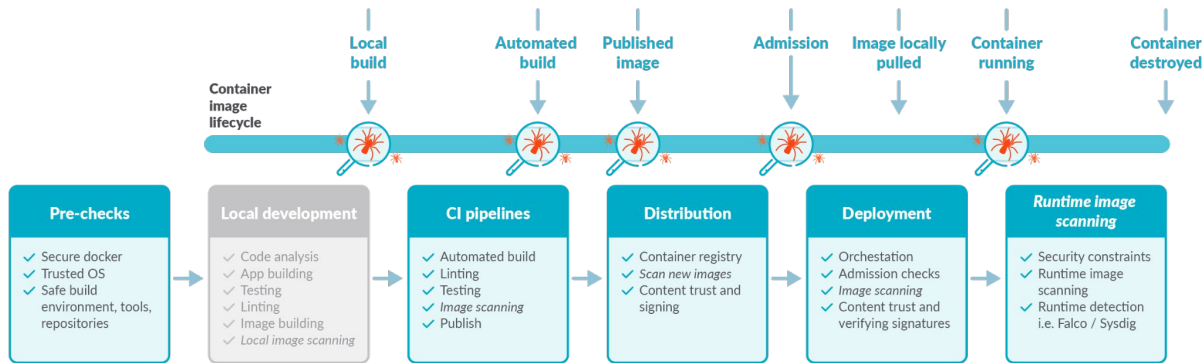
# Image Scanning

- **Image scanning** can be implemented at different stages
  - Detect bad practices and known vulnerabilities
  - The earlier the scan is performed, the better



# Image Scanning





## 5. Beyond image building

# Beyond image building

So far we talked about "making" coffee...



# Beyond image building

Now it is time to enjoy our cup of coffee



# Beyond image building - runtime

- Remember **unnecessary privileges**?
  - The **orchestrator** or **runtime** environment (i.e., docker run, kubernetes, etc.) has the last word on who is the running container effective user.
  - Avoid **running your environment as root**
  - Openshift and some Kubernetes clusters will apply **restrictive policies by default**, **preventing root** containers from running or using a **random UID**

# Beyond image building - runtime

- **Restrict application capabilities** on runtime
  - In case your container is compromised, the range of action available to an attacker is limited
  - [--cap-drop flag](#) in Docker
  - [securityContext.capabilities.drop](#) in Kubernetes
  - AppArmor in [Docker](#) or [Kubernetes](#)
  - Seccomp in [Docker](#) or [Kubernetes](#).

# Beyond image building - runtime detection


- New vulnerabilities are **discovered daily**
- Remember: Update often + versioning strategy
- Runtime detection is key
  - **Re-evaluate** scanned images to detect new applying vulnerabilities
  - **Runtime threat detection**. Falco can help.
  - **Respond** to suspicious activity. i.e. stop or pause container + **forensics**





# Beyond image building - mutant tags

- **Beware!** *Attack of the mutant tags*



<https://sysdig.com/blog/toctou-tag-mutability/>

<https://www.youtube.com/watch?v=j8K6EjOPhxs>

# Beyond image building - and more!

- The docker socket is [a big privileged door into your host system](#)
  - Make sure your `/var/run/docker.sock` has the correct permissions
  - If docker is exposed via TCP, make sure it is properly protected.

# Beyond image building - and more!

- The docker socket is a big privileged door into your host system
  - Make sure your `/var/run/docker.sock` has the correct permissions
  - If docker is exposed via TCP, make sure it is properly protected.
- Use [docker content trust](#), Docker notary, Harbor notary, or similar tools to digitally sign your images and then verify them on runtime.

# Beyond image building - and more!

- The docker socket is a big privileged door into your host system
  - Make sure your `/var/run/docker.sock` has the correct permissions
  - If docker is exposed via TCP, make sure it is properly protected.
- Use docker content trust, Docker notary, Harbor notary, or similar tools to digitally sign your images and then verify them on runtime.
- Use Docker health-checks or Kubernetes livenessProbes
  - Critical for long running or persistent services in order to ensure they are healthy



# Q & A



**sysdig**

Seeing is **Securing**